

# STA2201H Methods of Applied Statistics II

Monica Alexander

Week 5: Bayesian regression and Stan

# Announcements

- ▶ Assignment 1 being graded
- ▶ Assignment 2 coming soon

## Where are we at

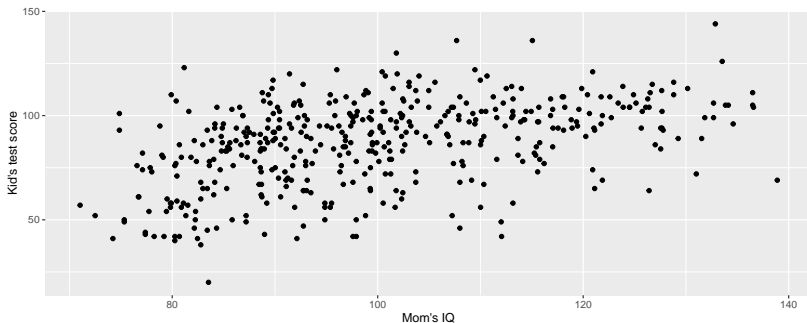
- ▶ Bayesian inference revolves around inference based on the posterior
- ▶ Posterior usually hard to write down in closed form
- ▶ But as long as we can get a set of samples from posterior, we can do inference
- ▶ For most problems, we can construct an MCMC algorithm that can be used to generate samples from posterior distributions
- ▶ lots of standard software to run MCMC so that we (usually) don't have to code it ourselves
- ▶ We will be using Stan, which fits models using a version of HMC

## Bayesian inference for regression models

# Kid's scores

- ▶ Outcome is Kid's test scores
- ▶ Let's introduce a covariate/explanatory variable of Mom's IQ

1) Question / goal : Describe the association between kid's test scores and Mom's IQ



# Scientific model

- 2) What is the Scientific model (how are these observed data generated?)

How does Mom's IQ influence Kid's score? If we think about this relationship causally

$$X \rightarrow Y$$

- ▶ Changing Mom's IQ would change Kid's test score, but not the other way around
- ▶ This is a scientific claim

# Scientific model

Adding another piece to our scientific model

$$X \rightarrow Y \leftarrow U$$

“Kid’s score is a function of Mom’s IQ and other stuff” This implies we need to find some function  $Y = f(X, U)$ . Let’s assume Kid’s score is a proportion of Mom’s IQ plus the influence of unobserved causes

## Statistical model

A reasonable model to consider is

$$y_i | \mu_i, \sigma \sim N(\mu, \sigma^2)$$

$$\mu_i = \alpha + \beta x_i$$

where  $X_i$  is mother's IQ score. This is a simple linear regression model. We are primarily interested in obtaining estimates for the regression coefficients,  $\alpha$  and  $\beta$ .

We need to put priors on  $\sigma$  (as before) but also  $\alpha$  and  $\beta$ . Let's put

$$\alpha \sim N(0, 100^2)$$

$$\beta \sim N(0, 10^2)$$

$$\sigma \sim \text{Half-Normal}(0, 1)$$



# Bayesian regression

- ▶ OLS or MLE finds estimates of the parameters that best fit the data
- ▶ Bayesian inference incorporates prior information about the parameters
- ▶ In Bayesian inference, the estimates are a compromise between the prior info and the data

# Bayesian inference for linear regression

What does Bayesian inference get us that MLE doesn't?

- ▶ **Inclusion of prior information:**

- ▶ we usually know something
- ▶ makes inferences more stable, as the estimates are typically somewhere between the prior and what would be obtained from the data alone

- ▶ **Propagation of uncertainty:**

- ▶ least squares gives us a point estimate
- ▶ in Bayesian inference, we can summarize uncertainty using simulations from the posterior distribution

## Posterior distribution

$$\Pr(\alpha, \beta, \sigma \mid Y_i, X_i) = \frac{\Pr(Y_i \mid X_i, \alpha, \beta, \sigma) \Pr(\alpha, \beta, \sigma)}{Z}$$

- Note that  $\alpha$  and  $\beta$  describe the line (conditional expectation) and  $\sigma$  describes the variation around the line

## Prior predictive distributions

- ▶ Priors should express scientific knowledge, but “softly”
- ▶ Sigma must be positive
- ▶ Kid score on average increases with Mom IQ?
- ▶ ???

Idea of prior predictive distributions:

- ▶ We can understand the implications of priors through simulation: check that before the model sees data, it doesn't hallucinate impossible things.
- ▶ We can force the model to make predictions even before data.

## Prior predictive distributions

- ▶ If we specify proper priors for all parameters in the model, our model is **generative**
- ▶ Yields a joint prior distribution on the parameters and data, and hence a prior marginal distribution for the data

Prior predictive distribution for new  $\tilde{y}$


$$p(\tilde{y}) = \int_{\Theta} p(\tilde{y}, \theta) d\theta = \int_{\Theta} p(\tilde{y}|\theta) p(\theta) d\theta$$

In practice (in R) we can simulate values of  $\theta$  from the prior distribution(s), and then simulate from the likelihood to generate values of  $\tilde{y}$ , and then look at the resulting distribution.

For now, I'm just going to generate values of the conditional expectation/linear predictor.

# Make some lines

```
n <- 1000  
alpha <- rnorm(n, 0, 100)  
beta <- rnorm(n, 0, 10)  
plot(NULL, xlim=c(20, 150), ylim = c(20, 150), xlab = "Mom IQ", ylab = "Kid score")  
for (j in 1:50) abline(a = alpha[j], b = beta[j], col = 2, lwd = 2)
```



## Sermon on priors (from Stat Rethinking)

- ▶ There are no correct priors, only scientifically justifiable priors
- ▶ Justify with information outside the data, like the rest of the model (eg the generative model)
- ▶ Priors are not so important in simple models
- ▶ Very important/useful in complex models
- ▶ Need to simulate and understand behavior

# In Stan

```
data {  
  int<lower=0> N;           // number of kids  
  int<lower=0> K;           // number of covariates  
  vector[N] y;             // scores  
  matrix[N, K] X;          // design matrix  
}  
  
parameters {  
  real alpha;  
  vector[K] beta;  
  real<lower=0> sigma;  
}  
  
transformed parameters {  
}  
  
model {  
  //priors  
  alpha ~ normal(0, 100);  
  beta ~ normal(0, 1);  
  sigma ~ normal(0,1);  
  
  //likelihood  
  y ~ normal(alpha + X*beta, sigma);  
}
```



# Fits comparison


```
summary(fit)$summary[c("alpha", "beta[1]"),]
```

```
##              mean      se_mean      sd      2.5%      25%      50%
## alpha  25.9021707 0.167995802 5.90152220 14.4842607 21.9706545 25.9697777
## beta[1] 0.6087277 0.001649548 0.05814357 0.4900049 0.5704404 0.6081894
##              75%      97.5%    n_eff    Rhat
## alpha  29.8030582 37.6457406 1234.046 1.004106
## beta[1] 0.6476124 0.7212384 1242.434 1.004011
```

```
summary(lm(kid_score~mom_iq, data = kidiq))
```

```
##
## Call:
## lm(formula = kid_score ~ mom_iq, data = kidiq)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -56.753 -12.074   2.217  11.710  47.691
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  25.79978     5.91741    4.36 1.63e-05 ***
## mom_iq        0.60997     0.05852   10.42 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 18.27 on 432 degrees of freedom
## Multiple R-squared:  0.201, Adjusted R-squared:  0.1991
## F-statistic: 108.6 on 1 and 432 DF, p-value: < 2.2e-16
```

```
pairs(fit, pars = c("alpha", "beta[1]"))
```



## What do we get

```
post_samples <- extract(fit)
length(post_samples)
```

```
## [1] 4
```

```
names(post_samples)
```

```
## [1] "alpha" "beta"  "sigma" "lp__"
```

## What do we get

```
dim(post_samples[["alpha"]])
```

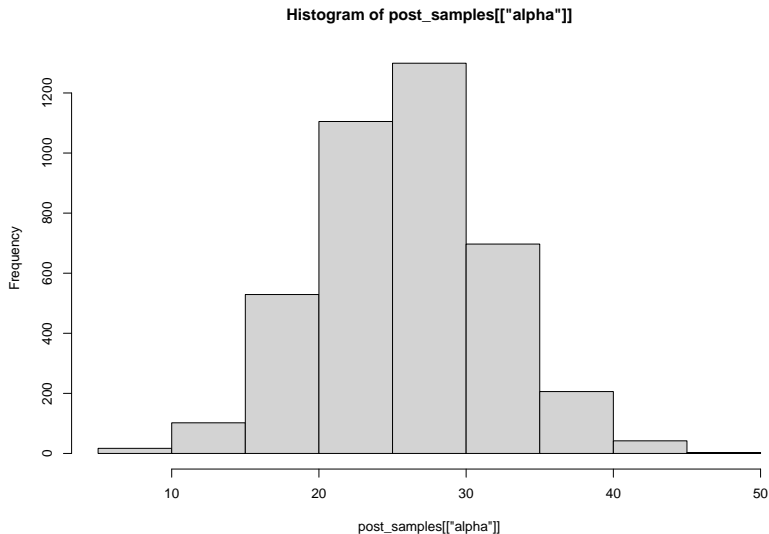
```
## [1] 4000
```

```
post_samples[["alpha"]][1:5]
```

```
## [1] 31.40069 27.90220 30.84657 25.99818 25.61134
```

# What do we get

```
hist(post_samples[["alpha"]])
```



# Tidy version

```
library(tidybayes)
fit |>
  gather_draws(alpha)
```

```
## # A tibble: 4,000 x 5
## # Groups:   .variable [1]
##   .chain .iteration .draw .variable .value
##   <int>     <int> <int> <chr>    <dbl>
## 1       1         1     1 1 alpha    29.0
## 2       1         2     2 2 alpha    34.6
## 3       1         3     3 3 alpha    24.8
## 4       1         4     4 4 alpha    29.2
## 5       1         5     5 5 alpha    25.7
## 6       1         6     6 6 alpha    21.0
## 7       1         7     7 7 alpha    20.8
## 8       1         8     8 8 alpha    20.1
## 9       1         9     9 9 alpha    29.8
## 10      1        10    10 10 alpha    17.6
## # i 3,990 more rows
```

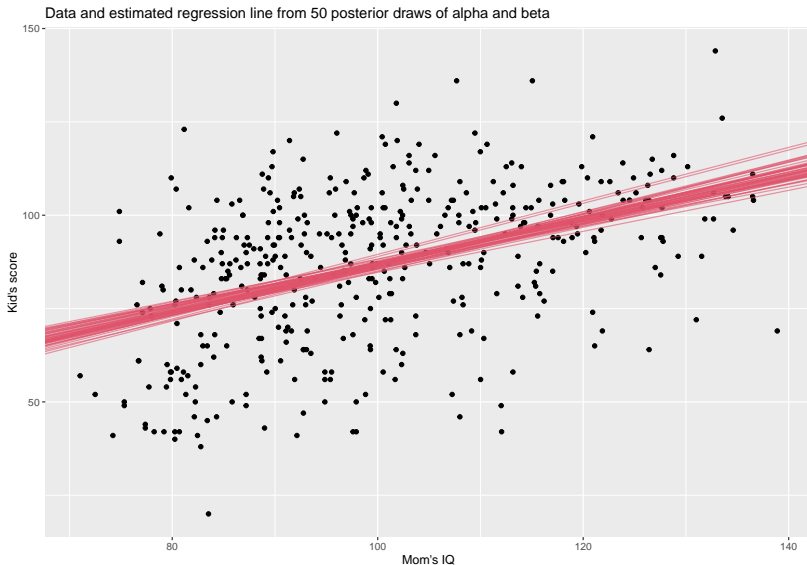
# What can we do

- ▶ The data and model are combined to form a posterior distribution, which we typically summarize by a set of simulations of the parameters in the model
- ▶ We can propagate uncertainty in this distribution, that is, we can get simulation-based prediction for unobserved or future outcomes that accounts for uncertainty in the model parameters

With simulations, we can

- ▶ Visualize uncertainty in the regression line
- ▶ Get uncertainty for functions of parameters
- ▶ Make predictions based on new data points


# The posterior is full of lines





# Uncertainty about a function of parameters

For example, posterior samples for the ratio of  $\alpha$  and  $\beta$



# Making predictions

Consider making a prediction of kid's score with a new observation of mother's IQ,  $x^{\text{new}}$ . We have

- ▶ the point prediction  $\hat{\alpha} + \hat{\beta}x^{\text{new}}$
- ▶ the linear predictor with uncertainty  $\alpha + \beta x^{\text{new}}$ 
  - ▶ propagates uncertainty in regression coefficients
  - ▶ represents the distribution of uncertainty about the expected value of  $y$  for new data points  $x^{\text{new}}$
- ▶ the predictive distribution for a new observation  $\alpha + \beta x^{\text{new}} + \text{error}$ 
  - ▶ represents uncertainty about a new observation  $y$  with predictor  $x^{\text{new}}$

# Predictions

Consider a new mother with an IQ of 110.

Point prediction: use medians of posterior samples for  $\hat{\alpha}$  and  $\hat{\beta}$

```
x_new <- 110
alpha_hat <- median(post_samples[["alpha"]])
beta_hat <- median(post_samples[["beta"]])


alpha_hat + beta_hat*x_new
```

```
## [1] 92.87062
```

# Predictions

Linear predictor with uncertainty:


```
alpha <- post_samples[["alpha"]]  
beta <- post_samples[["beta"]][,1]  
  
lin_pred <- alpha + beta*x_new  
hist(lin_pred)
```



# Predictions

Predictive distribution for new observation:

```
sigma <- post_samples[["sigma"]]  
y_new <- rnorm(n = length(sigma), mean = lin_pred, sd = sigma)  
hist(y_new)
```



## Can also do this within Stan

Can get posterior predictive distribution samples using the generated quantities block:

```
generated quantities{  
  real y_new[1];  
  y_new = normal_rng(alpha + x_new*beta, sigma);  
}
```

## Posterior predictive distribution

$$p(\tilde{y}|y) = \int_{\Theta} p(\tilde{y}|\theta, y)p(\theta|y)d\theta$$

- ▶ After we have seen the data and obtained the posterior distributions of the parameters, we can now use the posterior distributions to generate new data from the model.
- ▶ Given the posterior distributions of the parameters of the model, the posterior predictive distribution gives us some indication of what new data might look like, given the data and model.
- ▶ We can avoid performing the integration explicitly by generating samples from the posterior predictive distribution.

Posterior predictive distributions also important for model checking. More next week.

# Posterior predictive distribution

Posterior predictive distribution for new  $\tilde{y}$

$$p(\tilde{y}|y) = \int_{\Theta} p(\tilde{y}|\theta, y)p(\theta|y)d\theta$$

To obtain samples from this distribution, we need to


- ▶ Get posterior samples of our parameters  $\theta^{(s)}$  (MCMC output!)
- ▶ For each posterior sample, we obtain one replicated dataset  $\tilde{y}^{(s)}$  by sampling from the likelihood  $p(\tilde{y}|\theta^{(s)})$ . Can do this in R or within Stan.



Centering predictors to improve posterior geometries

# Remember this

```
pairs(fit, pars = c("alpha", "beta"))
```



# Centering

```
data <- list(y = y,  
            N = length(y),  
            K = 1,  
            X = as.matrix(kidiq$mom_iq - mean(kidiq$mom_iq)))  
fit2 <- stan(file = "kids3.stan",  
            data = data)
```

# Summary of fit

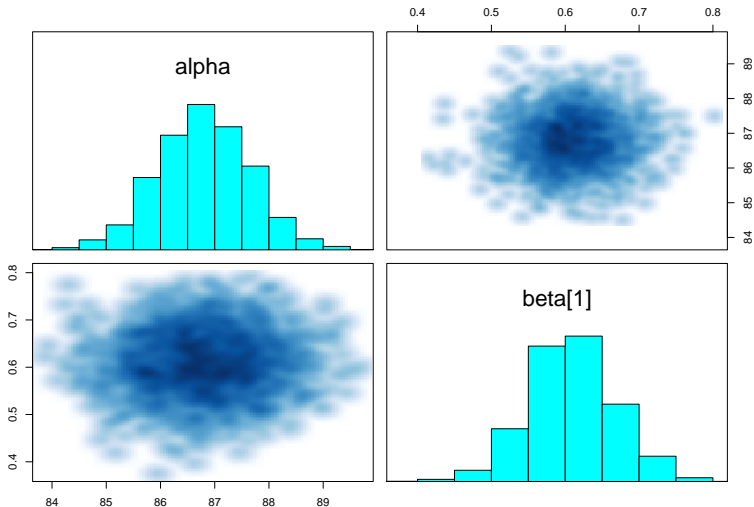
```
summary(fit2)$summary[c("alpha", "beta[1]"),]
```

##		mean	se_mean	sd	2.5%	25%
##	alpha	86.8044586	0.0134982924	0.85488732	85.1236984	86.2366701
##	beta[1]	0.6094195	0.0008783611	0.05806288	0.4934787	0.5730991
##		75%	97.5%	n_eff	Rhat	
##	alpha	87.3905097	88.4824846	4011.069	0.9994836	
##	beta[1]	0.6474439	0.7245516	4369.694	0.9997044	

What's different? What's the same?

## Now look at joint posteriors

```
pairs(fit2, pars = c("alpha", "beta"))
```



What do you notice? Why does this matter?

## Centering predictors

- ▶ When the mean of the predictors is far away from zero, changes in the slope induce the opposite change in the intercept
- ▶ Hard to interpret what intercepts mean
- ▶ Harder to sample: reducing correlation may speed up convergence

Changing prior information

## Changing prior information

What if we knew with relative certainty that there's a 1:1 correspondence between kid's score and mother's IQ? How would we encode this information?



# Changing prior information

$$\beta \sim N(1, 0.01^2)$$

Let's fit this:


```
data <- list(y = y,  
             N = length(y),  
             K = 1,  
             X = as.matrix(kidiq$mom_iq - mean(kidiq$mom_iq)))  
fit3 <- stan(file = "kids5.stan",  
             data = data)
```

# Summary of fit

```
summary(fit3)$summary[c("alpha", "beta[1]"),]
```

##	mean	se_mean	sd	2.5%	25%	50%
## alpha	86.7887933	0.0124096015	0.762758716	85.2809431	86.2911873	86.7894400
## beta[1]	0.9848027	0.0001461138	0.009825454	0.9658015	0.9781669	0.9848722
##	75%	97.5%	n_eff	Rhat		
## alpha	87.3024986	88.255837	3777.972	1.0000040		
## beta[1]	0.9914099	1.004173	4521.919	0.9999344		

# Comparison with weakly informative priors



# Comments

- ▶ Okay, maybe this was a bad decision in this context, but when might we want to consider more informative priors?
- ▶ Measurement error?
- ▶ Less data?
- ▶ Previous evidence?

# Break the model


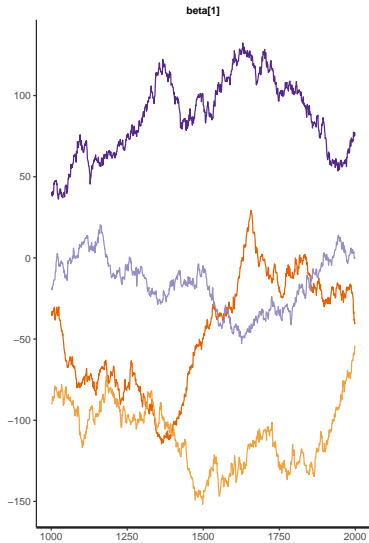
$$y_i | \mu_i, \sigma \sim N(\mu, \sigma^2)$$

$$\mu_i = \alpha + \beta_1 x_i + \beta_2 x_i$$

Priors on  $\beta$  are improper:  $p(\beta) \propto 1$

```
data <- list(y = y,  
            N = length(y),  
            K = 2,  
            X = cbind(as.matrix(kidiq$mom_iq), as.matrix(kidiq$mom_iq)))  
fit4 <- stan(file = "kid6.stan",  
            data = data)
```

```
## Inference for Stan model: kid6.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##               mean se_mean      sd      2.5%      25%      50%      75%      97.5%
## alpha         26.51    0.92  4.74    17.83    23.17    26.27    29.49    36.84
## beta[1]      -21.07   49.83 75.08   -134.71   -86.50   -22.52    30.97   118.41
## beta[2]       21.67   49.83 75.08   -117.87   -30.42    23.14    87.10   135.35
## sigma        14.86    0.10  0.39    14.07    14.63    14.81    15.06    15.72
## lp__        -1606.96    0.19  1.33  -1610.24  -1607.70  -1606.59  -1605.91  -1605.42
##               n_eff Rhat
## alpha         26 1.16
## beta[1]        2 4.39
## beta[2]        2 4.39
## sigma         15 1.43
## lp__          48 1.08
##
## Samples were drawn using NUTS(diag_e) at Wed Feb  7 08:50:52 2024.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```



# Compare to weakly informative priors

Priors on  $\beta$  are  $\beta \sim N(0, 1)$

```
data <- list(y = y,  
            N = length(y),  
            K = 2,  
            X = cbind(as.matrix(kidiq$mom_iq), as.matrix(kidiq$mom_iq)))  
fit5 <- stan(file = "kids3.stan",  
            data = data)
```

What do you think will happen?




# Results

What is identifiable given the observed data?

```
## Inference for Stan model: kids3.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##               mean se_mean   sd      2.5%      25%      50%      75%      97.5%
## alpha         25.20    0.33  5.77     13.68     21.25     25.24     29.06     36.50
## beta[1]        0.22    0.66  6.95    -13.44     -4.55      0.36      4.88     14.02
## beta[2]        0.39    0.66  6.95    -13.40     -4.26      0.24      5.17     14.09
## sigma         18.28    0.04  0.61      17.11     17.85     18.27     18.70     19.52
## lp__          -1477.54    0.06  1.39   -1481.19   -1478.27   -1477.22   -1476.47   -1475.82
##               n_eff Rhat
## alpha         308 1.01
## beta[1]       111 1.01
## beta[2]       111 1.01
## sigma        222 1.01
## lp__          472 1.01
##
## Samples were drawn using NUTS(diag_e) at Wed Feb  7 08:52:46 2024.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

# Traceplots



Shrinkage priors

## Additional reading for this part

- ▶ Piironen and Vehtari, 2017. 'Sparsity information and regularization in the horseshoe and other shrinkage priors'. Electron. J. Statist. 11(2): 5018-5051 (2017). DOI: 10.1214/17-EJS1337SI
- ▶ Nice but long case study by Michael Betancourt: [https://betanalpha.github.io/assets/case\\_studies/modeling\\_sparsity.html](https://betanalpha.github.io/assets/case_studies/modeling_sparsity.html)

## Penalized regression models

- ▶ In many contexts, may have lots of possible covariates ( $p$ ), which is big in relation to the number of observations ( $n$ )
- ▶ Common problem in genomic studies, imaging, etc
- ▶ In general, when there are many possible covariates, overfitting can be a problem
- ▶ (You may have seen) LASSO/Ridge regression (in frequentist-based inference), which maximize an objective function defined as the log-likelihood minus a penalty term.

# Penalization in Bayesian models

- ▶ In a Bayesian model, we can use induce sparsity through priors on regression coefficients, which assume only a small number of covariates are non-zero.
- ▶ One way of doing this is to choose a prior on the coefficient that has a sharp peak at 0, but also has a heavy tail, allowing some coefficient estimates to 'escape'

## Horseshoe prior

The horseshoe prior (Carvalho, Polson, and Scott 2009) sets the scale for each component to the product of a global scale,  $\tau$  and a local scale,  $\lambda_j$  which are themselves unknown parameters:

$$\beta_j | \lambda_j, \tau \sim N(0, \tau^2 \lambda_j^2)$$


$$\lambda_j \sim C^+(0, 1)$$

$$\tau \sim C^+(0, \tau_0)$$

- ▶ where  $C^+$  is the half-Cauchy distribution
- ▶ Concept of **global** and **local** scales: the global scale ( $\tau$ ) shrinks all coefficients to zero, while the local scale ( $\lambda_j$ ) allows some coefficients to escape shrinkage
- ▶ Different levels of sparsity can be achieved by changing the value of  $\tau$

## Horseshoe prior notes

- ▶ This is a continuous version of the spike-and-slab prior (Mitchell and Beauchamp, 1988; George and McCulloch, 1993)
- ▶ It has a **hierarchical** structure, in that the local scales are assumed to be exchangeable (shelve this, more later)
- ▶ Compared to a normal prior, horseshoe prior has more density at 0, but also more density for extreme values.
- ▶ Thus, for coefficients with very weak evidence, the regularizing prior will shrink it to zero, whereas for coefficients with strong evidence, the shrinkage will be very small.






## Why horseshoe?

Given the hyperparameters, the posterior mean satisfies approximately


$$\bar{\beta}_j = (1 - \kappa_j) \beta_j^{\text{ML}}, \quad \kappa_j = \frac{1}{1 + n\sigma^{-2}\tau^2\lambda_j^2}$$

where  $\kappa_j$  is the shrinkage factor. With  $\lambda_j \sim C^+(0, 1)$  the prior on  $\kappa_j$  looks like a U. Below is when  $n\sigma^{-2}\tau^2 = 0.9$



## Changing $\tau$

When  $n\sigma^{-2}\tau^2 = 0.1$  (i.e. small  $\tau$ ), there's more shrinkage of the coefficients to zero.



# In Stan

```
int<lower=0> N;           // number of kids
int<lower=0> K;           // number of covariates
vector[N] y;            // scores
matrix[N, K] X;         // design matrix
}
parameters {
  real alpha;
  vector[K] beta;
  real<lower=0> sigma;
  real<lower=0> tau;
  real<lower=0> lambda[K];
}
transformed parameters {
}
model {
  //priors
  alpha ~ normal(0, 100);
  for(k in 1:K){
    beta[k] ~ normal(0, tau*lambda[k]);
  }
  sigma ~ normal(0,10);
  lambda ~ cauchy(0,1);
  tau ~ cauchy(0,1);

  //likelihood
  y ~ normal(alpha + X*beta, sigma);
}
```

## Regularized horseshoe

- ▶ Horseshoe prior does not regularize slopes that are far from zero at all
- ▶ The regularized horseshoe (or Finnish horseshoe, or pony horseshoe???) introduces an additional layer to further regularize larger coefficients:

$$\beta_j | \lambda_j, \tau \sim N(0, \tau^2 \tilde{\lambda}_j^2)$$

$$\tilde{\lambda}_j = \frac{c \lambda_j}{\sqrt{c^2 + \tau^2 \lambda_j^2}}$$

$$\lambda_j \sim C^+(0, 1)$$

$$c^2 \sim \text{Inv-Gamma}(\nu/2, \nu/2s^2)$$

$$\tau \sim C^+(0, \tau_0)$$

## Regularized horseshoe

$$\beta_j | \lambda_j, \tau \sim N(0, \tau^2 \tilde{\lambda}_j^2)$$

$$\tilde{\lambda}_j = \frac{c \lambda_j}{\sqrt{c^2 + \tau^2 \lambda_j^2}}$$

$$\lambda_j \sim C^+(0, 1)$$

$$c^2 \sim \text{Inv-Gamma}(\nu/2, \nu/2s^2)$$

$$\tau \sim C^+(0, \tau_0)$$

- ▶ New scale variable  $c$ , which controls the distribution of coefficients far from zero (“slab width”)
- ▶ When  $\tau^2 \lambda_j^2 \ll c^2$ , meaning the coefficient is close to zero, then  $\tilde{\lambda}^2 \rightarrow \lambda^2$
- ▶ When  $\tau^2 \lambda_j^2 \gg c^2$  then  $\tilde{\lambda}^2 \rightarrow c^2/\tau^2$  and the prior on the coefficient approaches  $N(0, c^2)$ .


# Regularized horseshoe

- ▶ Often performs better computationally than normal horseshoe (helps to reduce the number of divergences)
- ▶ Be wary of sensitivities to hyperparameter choice (i.e.  $\tau_0$ ,  $\nu$  and  $s$ )
- ▶ Example Stan code is up on github

## Kid's scores with all possible covariates and interactions


- ▶ Covariates are mother's education (high school y/n), age, IQ, and work status
- ▶ Included all interactions
- ▶ NOTE: regularization treats variables which vary on a larger scale as more relevant, so should scale variables such that all have unit variance before fitting model

# Standard regression






# Horseshoe



# Regularized horseshoe



# Summary

In any modeling problem:

- 1) Question/Goal
- 2) Scientific model
- 3) Statistical model

Bayesian inference for linear regression

- ▶ Focus on simulation-based inference and prediction, rather than point estimates
- ▶ Can simulate predictions even before seeing data
- ▶ Easy to propagate uncertainty to predictions, functions of estimated parameters

# Summary

## Inducing sparsity

- ▶ Rather than inducing sparsity through a penalty term, we induce sparsity through priors
- ▶ Horseshoe priors shrink everything to zero, with a local scale variable allowing some coefficients to escape shrinkage
- ▶ Regularized horseshoe controls the scale of non-zero coefficients, which is often better computationally

Lab: practice with kids dataset